

The slide or note you've shared highlights **three major complexities in stream processing** — data, systems, and workloads. Let's break each one down with context and examples:

1. COMPLEX DATA

✓ Challenges:

- **Diverse data formats:** Data might arrive in different formats like:
 - **JSON** – readable structured data
 - **Avro** – compact binary format for serialization
 - **Binary** – unreadable raw data that needs decoding
- **Dirty data:** Incoming streams may have:
 - Missing fields, corrupt records, or wrong types (e.g., timestamp is a string instead of a number)
- **Late or out-of-order data:**
 - Events might not arrive in the order they were generated.
 - For example, a sensor might send data with timestamps like 10:00, 10:02, 10:01.

🔧 What this means for stream processing:

You need tools that can **parse multiple formats**, **clean or validate data in real-time**, and **handle event-time vs. processing-time** (using watermarks or windowing logic).

2. COMPLEX SYSTEMS

✓ Challenges:

- **Diverse storage systems:**

- A stream processor might pull data from **Kafka**, but enrich it using **RDBMS**, store results in **S3**, and send alerts via **Kinesis**.
- **System failures:**
 - Streams need to be **fault-tolerant** — what happens if Kafka goes down? Or your job crashes?
 - You must support **exactly-once or at-least-once** semantics depending on your use case.

What this means:

Your processing engine (like **Flink**, **Spark Streaming**, or **Kafka Streams**) needs to be able to **connect to various sources/sinks, recover from crashes, and maintain consistency**.

3. COMPLEX WORKLOADS

Challenges:

- **Streaming + interactive queries:**
 - Users might want dashboards that run **real-time + historical** queries together.
 - For example, show today's order rate **vs.** average order rate over the last week.
- **Machine learning:**
 - Stream processing might involve:
 - Real-time predictions (e.g., fraud detection)
 - Online learning (models updated continuously)
 - Model inference in the stream

What this means:

You need systems that allow **low-latency queries**, support **caching**, and can **integrate ML pipelines** either via APIs or embedding models into the stream logic.

Summary Table:

Complexity	Challenges	Needs
Complex Data	Dirty, out-of-order, multi-format data	Schema handling, cleansing, windowing
Complex Systems	Many storage layers, system failures	Integration, checkpointing, exactly-once semantics
Complex Workloads	Real-time + batch, ML inference	Low-latency queries, scalable ML integration

Would you like a visual diagram or example pipeline showing how these come together in a real system (e.g., fraud detection or real-time analytics)?

This slide titled "**Anatomy of a Streaming Query**" explains the **source** part of a streaming pipeline in **Apache Spark Structured Streaming**.

◆ Code Breakdown:

python

CopyEdit

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
```

✓ What it does:

- **spark.readStream**: Tells Spark to read data as a **streaming source**, rather than a static DataFrame.
- **.format("kafka")**: Specifies that the source is Kafka.
- **.option("subscribe", "input")**: Subscribes to a Kafka topic named **"input"**.
- **.load()**: Actually loads the stream and returns a streaming DataFrame.

♦ Right-hand Side Explanation:

Source – Key Concepts:

- **"Specify one or more locations to read data from"**
You can stream from a single source or multiple (e.g., Kafka, files, sockets).
- **"Built-in support for Files/Kafka/Socket, pluggable."**
Spark comes with built-in connectors for:
 - **Kafka** (common for real-time logs, events)
 - **Files** (like CSV or JSON in a folder being populated over time)
 - **Socket** (like streaming text from a TCP server)
 - You can plug in other custom sources (e.g., MQTT, Kinesis, Pub/Sub).

"Can include multiple sources of different types using `union()`"

If needed, you can combine streams from multiple sources (like Kafka + file-based ingestion) using:

```
python
CopyEdit
df1.union(df2)
```

-

🧠 Real-world Use Case:

Imagine you're building a real-time analytics dashboard:

- You want to **read events from Kafka** (user clicks).
- You also want to **read config files** from S3 that update periodically.
- Spark lets you read both as streaming sources and **merge them** with `union()` to power downstream analytics.

Code Breakdown:

```
python
CopyEdit
spark.readStream
    .format("kafka")
    .option("subscribe", "input")
    .load()
    .groupBy('value.cast("string") as key')
    .agg(count("*") as 'value')
```

✓ What it does:

- **Reads from Kafka** like in the previous slide.
- **Casts the Kafka value to a string**, and **groups** the data by that value (acting as a key).
- **Aggregates** (counts) how many times each key appears in the stream.

This transformation helps perform **real-time analytics**, like counting how often a message appears in the stream.

♦ Right-hand Side Explanation:

Transformation – Key Concepts:

- **"Using DataFrames, Datasets and/or SQL"**
Spark allows you to write transformations using:
 - **DataFrame API**: like `.groupBy().agg()`
 - **Datasets** (typed APIs in Scala/Java)
 - **SQL**: by registering temporary views and running SQL queries.
- **"Catalyst figures out how to execute the transformation incrementally"**
Spark's **Catalyst optimizer**:
 - Automatically plans how to run the transformation.

- Ensures it only processes **new data (micro-batches)**, not redoing everything.
 - **"Internal processing always exactly-once"**
Spark guarantees **exactly-once semantics**:
 - Each event is processed **only once**, even during recovery from failure.
 - Checkpointing and write-ahead logs make this possible.
-

Real-world Use Case:

Let's say you're building a real-time **sentiment dashboard**:

- Users are sending reviews via Kafka.
- You want to **count how many reviews are positive, negative, or neutral** in real-time.
- You can group by the **sentiment** label and count them — exactly like this slide shows.

◆ **Code Block (Left Side):**

python

CopyEdit

```
input = spark.readStream
    .format("kafka")
    .option("subscribe", "topic")
    .load()

result = input \
    .select("device", "signal") \
    .where("signal > 15")

result.writeStream \
    .format("parquet") \
    .start("dest-path")
```

What this does:

- **Reads** streaming data from a Kafka topic.

- **Selects** the columns "device" and "signal".
 - **Filters** only the rows where signal > 15.
 - **Writes** the result to a Parquet file in dest-path.
-

◆ Center: Spark Planning Breakdown

1. Logical Plan

- This is what Spark generates from your **DataFrame operations**.
- It includes abstract steps:
 - **Read from Kafka**
 - **Project** (select device, signal)
 - **Filter** (signal > 15)
 - **Write to Kafka (or sink)**

2. Optimized Physical Plan

- Spark's **Catalyst optimizer** converts the logical plan into an efficient execution plan.
- Optimization includes:
 - **Operator fusion** (combining steps)
 - **Code generation** (for speed)
 - **Off-heap memory use**
- The plan includes:
 - Kafka as **source**
 - **Optimized operators**

- Kafka (or other) as **sink**
-

◆ Right: Incremental Execution

🔄 At each time step **t=1, t=2, t=3...**

- Spark **processes only the new data** arriving in the stream (called a **micro-batch**).
- It applies the same optimized plan on just the new input:
 - Filters rows
 - Projects columns
 - Writes results
- Spark creates a **new execution plan for each micro-batch**, enabling continuous and efficient streaming

◆ Code Breakdown (Left Side):

python

CopyEdit

```
spark.readStream
    .format("kafka")
    .option("subscribe", "input")
    .load()
    .groupBy('value.cast("string") as key')
    .agg(count("*") as 'value')
    .writeStream
    .format("kafka")
    .option("topic", "output")
```

✅ What it does:

- Reads from a Kafka topic called **"input"**.
- Groups and counts messages by their value.

- Writes the output to a **Kafka sink**, publishing results to a topic named "output".
-

♦ Right Side Explanation – Sink Concepts:

✓ What is a Sink?

A **sink** is the destination where **each processed micro-batch** is sent. Spark supports a variety of sinks like:

- **Kafka** – to publish back into a stream
 - **File systems** – to store in formats like CSV, Parquet, etc.
 - **Console** – for debugging
 - **Memory** – for in-app querying (not for production)
-

🧠 Key Points:

- **"Accepts the output of each batch."**
Spark processes streaming data in **micro-batches**, and each batch's result is sent to the sink.
 - **"Transactional and exactly once (when supported)."**
Some sinks like **files (e.g., Parquet)** and **Kafka (with idempotent writes)** support **exactly-once delivery semantics** to avoid duplicates or data loss.
 - **"Use foreach to execute arbitrary code."**
If you need to:
 - Send results to a **custom API**
 - Write to **external databases**
 - Trigger **notifications**
You can use `.foreach()` to run **custom code** for each row or batch.
-

💡 Example Use Cases for Sinks:

- Write aggregated metrics to a **dashboard via Kafka**
- Archive processed logs to **Amazon S3** using the file sink
- Store real-time fraud detection results into a **NoSQL DB** using `foreachBatch`

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as 'key)
  .agg(count("*") as 'value)
  .writeStream
  .format("kafka")
  .option("topic", "output")
  .trigger("1 minute")
  .outputMode("update")
```

Output mode – What's output

- Complete – Output the whole answer every time
- Update – Output changed rows
- Append – Output new rows only

} Trigger – When to output

- Specified as a time, eventually supports data size
- No trigger means as fast as possible

 databricks

```
spark.readStream
  .format("kafka")
  .option("subscribe", "input")
  .load()
  .groupBy('value.cast("string") as key')
  .agg(count("*") as 'value')
  .writeStream
  .format("kafka")
  .option("topic", "output")
  .trigger("1 minute")
  .outputMode("update")
  .option("checkpointLocation", "...")
  .start()
```

🧠 Real-World Example:

Let's say you're monitoring sensor data from IoT devices:

- You count how often a sensor reports a specific status (like “ALERT”).
- You group and aggregate per device and send that to Kafka every minute.

Without checkpointing:

- If the job fails at 10:03, all progress is lost.
- When restarted, it reprocesses from the beginning (duplicates results, re-counts data).

With checkpointing:

- If the job fails at 10:03, the last checkpoint (say at 10:02) allows it to resume from there.
- It ensures **exactly-once** output.

♦ Best Practices:

- Always set a **valid, persistent location** for "`checkpointLocation`" (e.g., `/mnt/checkpoints/iot-query`)
- Don't share a checkpoint directory across different queries
- Make sure the location is accessible to all nodes in the cluster

This slide dives deeper into **fault tolerance with checkpointing** in Spark Structured Streaming, explaining how it ensures **exactly-once guarantees** across failures or restarts.

♦ Core Concept: Checkpointing

Definition (from slide):

"Checkpointing – tracks progress (offsets) of consuming data from the source and intermediate state."

✅ What gets tracked?

- **Offsets:** The last successfully read position from a data source (e.g., Kafka offset).

- **State:** If you're doing stateful operations (like `groupBy + agg`), intermediate aggregation values are stored.
- **Metadata:** Info about execution plans, progress, schema, etc.

These are **persisted as JSON** in the **checkpoint directory** you define using:

```
.option("checkpointLocation", "/some/path/")
```

♦ Visual Explanation (Right side)

The timeline at **t = 1, 2, 3** shows:

- New data being processed in each micro-batch.
- After processing, **metadata and results are recorded into a write-ahead log**.
- In case of failure (e.g., crash after **t=2**), Spark:
 - Reloads the **checkpoint data**
 - Avoids reprocessing **t=1** and **t=2**
 - Resumes cleanly from **t=3**

This is how Spark achieves **end-to-end exactly-once guarantees**.

Real-World Example:

Imagine you're processing **credit card transactions** in real time to detect fraud.

Use case:

- Stream transaction data from Kafka.
- Use stateful logic: group by `card_id` and keep a running sum of spend in the last 10 minutes.

- Send alerts when suspicious thresholds are exceeded.

Without checkpointing:

- If Spark crashes, **state is lost**.
- It may miss alerts or double-send them (if it reprocesses old Kafka records).

With checkpointing:

- Spark saves:
 - Which transactions were already read (Kafka offsets)
 - Intermediate state (running sums)
- After a crash, Spark:
 - Picks up where it left off
 - Keeps your **fraud detection logic consistent**

♦ Key Takeaways:

Feature	Why It Matters
Offsets tracked	Prevents reprocessing of old data
State tracked	Ensures aggregations/window logic can continue
Metadata as JSON	Human-readable and easy to debug
Resilience after crash	Recovers query without duplicating output
Exactly-once guarantee	Critical for finance, alerts, analytics

Would you like to go over how **outputMode** or **trigger** settings affect checkpointing next?

```
val rawData = spark.readStream
  .format("kafka")
  .option("kafka.bootstrap.servers",...)
  .option("subscribe", "topic")
  .load()

val parsedData = rawData
  .selectExpr("cast (value as string) as json")
  .select(from_json("json", schema).as("data"))
  .select("data.*")

val query = parsedData.writeStream
  .option("checkpointLocation", "/checkpoint")
  .partitionBy("date")
  .format("parquet")
  .start("/parquetTable")
```



Great! You're referring to **Kafka source configuration options** used in **Spark Structured Streaming** via `.option(...)` when reading from Kafka.

Let's walk through this block **line by line**, explaining:

- How to use each option
- What it means
- What the syntax does

✓ How? (Basic connection to Kafka)

```
```scala
```

```
.option("kafka.bootstrap.servers", "broker1:9092,broker2:9092")
```

- **Required**: This is the **Kafka cluster** your stream will connect to.
- Provide a **comma-separated list** of Kafka brokers.
- This does **not** mean your app connects to all brokers — Kafka will handle that internally.

---

##  What? (How to specify topics or partitions)

### Option 1: **Fixed list of topics**

```
```scala
```

```
.option("subscribe", "topic1,topic2,topic3")
```

- Use when you **know exactly** which topics you want to read from.
- Spark reads **all partitions** from these topics.

Option 2: Dynamic topic pattern

```
.option("subscribePattern", "topic.*")
```

- Use a **regex pattern** to match topics dynamically.
- Spark will pick up any topic that matches `topic.*` — useful for logs, metrics, etc.
- **Topics created later** that match the pattern will be picked up too.

Option 3: Specific partitions (manual assignment)

```
.option("assign", """"{"topicA":[0,1]}""")
```

- JSON string specifying **which partitions to read** from which topic.
- Use this to control **exact partition-level consumption**.

- This **disables consumer group balancing** — you're fully in control.
-

✅ Where? (Offsets — where to start reading from)

Option: Set starting offset

```
.option("startingOffsets", "latest")
```

- Controls **where Spark starts reading from**:
 - **"latest"**: Only read **new messages** (messages produced **after** job starts).
 - **"earliest"**: Read **everything** from the beginning.

Advanced: Specific offsets per partition

```
.option("startingOffsets", """"{"topicA":{"0":23,"1":345}}""")
```

- Start reading from **offset 23 on partition 0** and **345 on partition 1** of topicA.
 - Useful for **resuming** from a specific point.
-

🧠 Example Code Putting It All Together

```
val df = spark.readStream
  .format("kafka")
  .option("kafka.bootstrap.servers", "localhost:9092")
  .option("subscribe", "logs,metrics")
  .option("startingOffsets", "earliest")
  .load()
```

✅ This connects to Kafka at **localhost:9092**, subscribes to **logs** and **metrics**, and starts reading from the **beginning**.

Summary Table

Option	Purpose	Value Type
<code>kafka.bootstrap.servers</code>	Connect to Kafka cluster	Comma-separated
<code>subscribe</code>	Read from fixed topics	Comma-separated
<code>subscribePattern</code>	Read from topics matching a regex	Regex string
<code>assign</code>	Read specific topic partitions	JSON string
<code>startingOffsets</code>	Where to begin reading messages	"latest", "earliest", or JSON

Let me know if you want to go over **endingOffsets**, checkpoint integration, or auto offset commits next!

Event time Aggregations

Windowing is just another type of grouping in Struct. Streaming

number of records every hour

```
parsedData
  .groupBy(window("timestamp", "1 hour"))
  .count()
```

avg signal strength of each device every 10 mins

```
parsedData
  .groupBy(
    "device",
    window("timestamp", "10 mins"))
  .avg("signal")
```

Support UDAFs!

Watermarking

Watermark - moving threshold of how late data is expected to be and when to drop old state

Trails behind **max seen event time**

Trailing gap is configurable



 databricks

This slide introduces **Arbitrary Stateful Operations** using Spark's `mapGroupsWithState` (introduced in **Spark 2.2**), which allows **fine-grained control over per-key state** in streaming applications.

Let's break it down with simple language and an example 📌

◆ What is `mapGroupsWithState`?

→ **A powerful Spark API for stateful stream processing**

It allows you to:

- Define your own **custom state** per key (like a counter, a timer, or any object)
 - Update, access, or remove this state as new data arrives
 - Set **timeouts per key** to clean up unused state
-

Syntax Overview from the Slide

```
ds.groupByKey(_._id)
  .mapGroupsWithState(timeoutConf)(mappingWithStateFunc)
```

Component	Role
<code>ds</code>	A dataset (structured streaming input)
<code>groupByKey(_._id)</code>	Group data by some key (e.g., user ID, sensor ID, etc.)
<code>mapGroupsWithState</code>	Apply a user-defined function that manages state per group
<code>te</code>	
<code>timeoutConf</code>	How long to retain state per key (event-time or processing-time)
<code>mappingWithStateFunc</code>	Your custom function that defines how to update and return output

◆ The Function Signature

```
def mappingWithStateFunc(
  key: K,
  values: Iterator[V],
  state: GroupState[S]
): U = {
  // Your logic here
}
```

Parameters:

- `key`: the current group key (e.g., a user ID)
- `values`: the new events for this key (batch of events)
- `state`: the **current state** for this key (you can read, update, or remove it)

Return:

- A result (**U**) for each group, like a transformed row or alert

Real-World Example: User Session Timeout

Say you're tracking user activity:

- Each message has **userId**, **action**, **timestamp**
- You want to maintain a **session window** per user
- If no activity from a user for 30 minutes, mark the session as closed

You can do this with **mapGroupsWithState**:

- Store last action timestamp as state
- If new event arrives after timeout → emit session close
- Otherwise → update session state

Key Features Highlighted in the Slide

Feature	What it means
mapGroupsWithState	Custom function per key with control over state
User-defined state	You choose what to store (count, last timestamp, etc.)
Per-key timeout	Clean up unused state after a configurable period
Event-time or processing-time	Timeouts can be based on event timestamp or system clock
Supports Scala and Java	Works with typed Datasets in both languages

Summary

- Use `mapGroupsWithState` when you need **custom logic** and **flexible state handling**.
- It's ideal for:
 - **Sessionization**
 - **Fraud detection**
 - **Last-known state tracking**
 - **Custom time-based alerts**

Let me know if you'd like to see a full Scala example with `GroupState` logic!

Great! This slide introduces **three powerful operations** in Spark Structured Streaming:

◆ 1. Streaming Deduplication

✓ Purpose:

To **remove duplicate records** from a stream — for example, when:

- Data has retry logic or network issues
- You don't want to process the same event twice

✓ Code:

```
parsedData.dropDuplicates("eventId")
```

- This keeps **only the first record** per unique `eventId`
- Internally, Spark stores a state of seen `eventIds` to avoid re-processing

⚠ Important:

You should combine this with **watermarking**:

```
parsedData
  .withWatermark("timestamp", "10 minutes")
  .dropDuplicates("eventId")
```

💡 This allows Spark to **forget old IDs after 10 minutes** and not hold memory forever.

♦ 2. Stream–Batch Joins

✅ Purpose:

Join **real-time data** with **static or slowly-changing data**, like:

- Enriching clickstream logs with user profile info
- Joining incoming orders with product catalog

✅ Code:

```
val batchData = spark.read
  .format("parquet")
  .load("/additional-data")
```

```
parsedData.join(batchData, "device")
```

- `batchData` is **loaded once** (not continuously)
- Spark joins every incoming streaming record with the batch table on `device`

⚠️ Tip:

The batch table should be small enough to **broadcast** to executors.

♦ 3. Stream–Stream Joins

✅ Purpose:

Join **two real-time data streams** together, like:

- Joining sensor readings with device status updates
- Matching orders with payments in real time

✓ Current Options:

- Use `mapGroupsWithState` for custom logic
- (Back when this slide was made) direct built-in support was *coming soon*
→ **Now it exists!** Spark supports **stream-stream joins** with watermark-based logic.

✓ Example (modern syntax):

```
stream1
  .withWatermark("ts1", "10 minutes")
  .join(
    stream2.withWatermark("ts2", "10 minutes"),
    expr("stream1.id = stream2.id AND ts1 BETWEEN ts2 - interval 5 minutes AND ts2 + interval 5 minutes")
  )
```

Summary Table

Operation	Description	Key Spark Features Used
Streaming Deduplication	Removes duplicates using a unique key	<code>dropDuplicates()</code> , Watermark
Stream–Batch Join	Enriches stream with static reference data	<code>join()</code>
Stream–Stream Join	Combines two streams using keys + time conditions	<code>join()</code> + Watermarks + Event-time windows

Let me know if you'd like to see examples of **handling late data in joins** or how **stream–batch joins** are optimized!